
Generative Adversarial Networks

GENERATE IMAGES BASED ON THE SMALLNORB DATASET

Dirk Stulgies, Prosper Kwabena Adjei

INTRODUCTION

- **This project was created as part of the `opencampus.sh` course **Generative Adversarial Networks**.**
- **The goal was to train a neural network with the `smallNorb` dataset and afterwards generate images similar to those in the original dataset.**

DATASET

- **Images of 50 toys.**
 - **Imaged by two cameras**
 - **Five categories** (four-legged animals, human figures, airplanes, trucks, and cars)
 - **Ten instances** (lion, hippo, van etc.)
 - **Six lighting conditions**
 - **Nine elevations** (30 to 70 degrees every 5 degrees)
 - **18 azimuths** (0 to 340 every 20 degrees)



III. 1: Images from the smallNorb dataset

NETWORK

- **Wasserstein GAN**
- **Gradient penalty**
- **Adam optimize**
- **BCE with logins loss function**
- **Initialise weights for Conv2D and Conv2DTranspose based on normal distribution.**

```
mean_iteration_critic_loss = 0
for _ in range(RUN_CRITIC_REPEAT):
    # Update the critic.
    crit_opt.zero_grad()
    fake_noise = get_noise(len(real), RUN_Z_DIM, device=RUN_DEVICE)
    fake = gen(fake_noise)
    crit_fake_pred = crit(fake.detach())
    crit_real_pred = crit(real)

    # Calculate the gradient penalty.
    epsilon = torch.rand(len(real), 1, 1, 1, device=RUN_DEVICE, requires_grad=True)
    gradient = get_gradient(crit, real, fake.detach(), epsilon)
    gp = gradient_penalty(gradient)
    crit_loss = get_crit_loss(crit_fake_pred, crit_real_pred, gp, RUN_CRITIC_LAMBDA)

    # Keep track of the average critic loss in this batch.
    mean_iteration_critic_loss += crit_loss.item() / RUN_CRITIC_REPEAT

    # Update the gradients.
    crit_loss.backward(retain_graph=True)

    # Update the optimizer.
    crit_opt.step()

# Add the critics loss value to the list, after each step.
crit_loss_step += [mean_iteration_critic_loss]
```

III. 2: Coded snippet, critics repeat, adding gradient penalty.

GENERATOR

- Five blocks containing:
 - ConvTranspose2d layer
 - InstanceNorm2d layer
 - LeakyReLU activation layer
 - Dropout layer
- The final block containing:
 - ConvTranspose2d layer
 - InstanceNorm2d layer
 - Tanh activation layer

```
self.layers = torch.nn.Sequential(  
    self.convTranspose2d_block(z_dim, hidden_dim, 3, 1, dropout_value, leak_value),  
    self.convTranspose2d_block(hidden_dim, hidden_dim, 3, 2, dropout_value, leak_value),  
    self.convTranspose2d_block(hidden_dim, hidden_dim, 5, 1, dropout_value, leak_value),  
    self.convTranspose2d_block(hidden_dim, hidden_dim, 5, 2, dropout_value, leak_value),  
    self.convTranspose2d_block(hidden_dim, hidden_dim, 7, 1, dropout_value, leak_value),  
    self.convTranspose2d_block(hidden_dim, 1, 7, 2, dropout_value, leak_value, final_layer=True)  
)
```

```
if not final_layer:  
    # Define the block for not final layer.  
    return torch.nn.Sequential(  
        torch.nn.ConvTranspose2d(input_channels, output_channels, kernel_size, stride),  
        torch.nn.InstanceNorm2d(output_channels),  
        torch.nn.LeakyReLU(leak_value, inplace=True),  
        torch.nn.Dropout(dropout_value)  
    )  
else:  
    # Define the block for the final layer.  
    return torch.nn.Sequential(  
        torch.nn.ConvTranspose2d(input_channels, output_channels, kernel_size, stride),  
        torch.nn.InstanceNorm2d(output_channels),  
        torch.nn.Tanh()  
    )
```

III. 3: Code snippet, generator structure.

CRITIC

➤ Five blocks containing:

- Conv2d layer
- InstanceNorm2d layer
- LeakyReLU activation layer
- Dropout layer

➤ The final block containing:

- Conv2d layer

```
self.layers = torch.nn.Sequential(  
    self.conv2d_block(img_channel, hidden_dim, 7, 2, dropout_value, leak_value),  
    self.conv2d_block(hidden_dim, hidden_dim, 7, 1, dropout_value, leak_value),  
    self.conv2d_block(hidden_dim, hidden_dim, 5, 2, dropout_value, leak_value),  
    self.conv2d_block(hidden_dim, hidden_dim, 5, 1, dropout_value, leak_value),  
    self.conv2d_block(hidden_dim, hidden_dim, 3, 2, dropout_value, leak_value),  
    self.conv2d_block(hidden_dim, 1, 3, 1, dropout_value, leak_value, final_layer=True)  
)
```


```
if not final_layer:  
    # Define the block for not final layer.  
    return torch.nn.Sequential(  
        torch.nn.Conv2d(input_channels, output_channels, kernel_size, stride),  
        torch.nn.InstanceNorm2d(output_channels),  
        torch.nn.LeakyReLU(leak_value, inplace=True),  
        torch.nn.Dropout(dropout_value)  
    )  
else:  
    # Define the block for the final layer.  
    return torch.nn.Sequential(  
        torch.nn.Conv2d(input_channels, output_channels, kernel_size, stride)  
    )
```

Ill. 4: Code snippet, critics structure.

RESULTS UNCONDITIONAL GAN



III. 5: Results after each epoch.



III. 6: Losses

REFERENCES

➤ Dataset

- smallNorb dataset: <https://cs.nyu.edu/~ylclab/data/norb-v1.0-small/>

➤ Couseira courses

- Build basic GAN's:
- Buils better GAN's
- Apply GAN's

➤ Code examples from other authors

- Reading the smallNorb files: <https://www.kaggle.com/code/leshabirukov/small-norb-load>
 - Generating images based on the smallNorb dataset: <https://medium.com/analytics-vidhya/applying-generative-adversarial-network-to-generate-novel-3d-images-ba70e1176dac>
-